

## Lab Course: Distributed Data Analytics Exercise Sheet 1

### Exercise 0

Processor	Intel core i5 7 <sup>th</sup> Gen
Operating system	Windows 10
Memory	8GB
Storage	Local Disk C: (299 GB)
Graphics	NVidia GeForce 920 MX
Programming language	Python

### Exercise 1

Solution file: mpimatvect.py

a)

Steps:

- 1) The rank is identified from the mpiexec command
- 2) Inputsize/rank is the batch size allocated to each of the processes except the last one.
- 3) The last process takes the remaining data after allocation to each process before it
- 4) Vector addition is performed in each slave
- 5) Results are returned to master

Vector addition

Core/input size	10 <sup>7</sup>	10 <sup>12</sup>	10 <sup>15</sup>
1	0.05500719975680113	-	-
2	1.803858700208366	-	-
3	1.844593099784106	-	-
4	1.8467364995740354	-	-

Note: For inputs of size 10<sup>12</sup>, 10<sup>15</sup>, the RAM is not big enough to perform operations of this magnitude

```
C:\Users\admin\Anaconda3\Scripts>mpiexec -n 1 python mpimatvect.py
rank 0
Traceback (most recent call last):
  File "mpimatvect.py", line 12, in <module>
    matrix_a = np.random.random((2,n))
  File "mtrand.pyx", line 430, in numpy.random.mtrand.RandomState.random
  File "mtrand.pyx", line 421, in numpy.random.mtrand.RandomState.random_sample
  File "_common.pyx", line 256, in numpy.random._common.double
MemoryError: Unable to allocate 14.6 TiB for an array with shape (2, 1000000000000) and data type float64
```

```
C:\Users\admin\Anaconda3\Scripts>mpiexec -n 1 python mpimatvect.py
rank 0
[1.18311933 1.5235067 0.97891462 0.96779836 1.13252036]
Time elapsed 0.00026950007304549217

C:\Users\admin\Anaconda3\Scripts>mpiexec -n 1 python mpimatvect.py
rank 0
[1.69407575 0.93302029 1.31820689 ... 1.04990507 0.257726 1.72507803]
Time elapsed 0.05500719975680113
```

## Exercise 1

Solution file – average\_vector.py

b)

Steps:

- 1) The rank is identified from the mpiexec command
- 2) Inputsize/rank is the batch size allocated to each of the processes except the last one.
- 3) The last process takes the remaining data after allocation to each process before it
- 4) Average is performed in each slave
- 5) Results are returned to master

Cores/Input size	10 <sup>7</sup>	10 <sup>12</sup>	10 <sup>15</sup>
1	0.012141900137066841	-	-
2	1.4384162002243102	-	-
3	1.4540423997677863	-	-
4	1.4918508999980986	-	-

Note: For inputs of size 10<sup>12</sup>, 10<sup>15</sup>, the RAM is not big enough to perform operations of this magnitude

```
rank 0
1000000000000000
Traceback (most recent call last):
  File "average_vector.py", line 11, in <module>
    matrix_a = np.random.random((1,n))
  File "mtrand.pyx", line 430, in numpy.random.mtrand.RandomState.random
  File "mtrand.pyx", line 421, in numpy.random.mtrand.RandomState.random_sample
  File "_common.pyx", line 256, in numpy.random._common.double_fill
MemoryError: Unable to allocate 7.28 TiB for an array with shape (1, 1000000000000000) and data type float64
```

```

C:\Users\admin\Anaconda3\Scripts>mpiexec -n 3 python average_vector.py
rank 1
10000000
matrix [[0.96702984 0.54723225 0.97268436 ... 0.0664702 0.11850422 0.77935834]]
result for rank 1 [1667208.27545623]
rank 2
10000000
matrix [[0.96702984 0.54723225 0.97268436 ... 0.0664702 0.11850422 0.77935834]]
result for rank 2 [1665801.67454719]
rank 0
10000000
matrix [[0.96702984 0.54723225 0.97268436 ... 0.0664702 0.11850422 0.77935834]]
result for rank 0 [1666430.26574922]
Averge of the vector [0.49994402]
Time 1.5075397002510726

```

Exercise 2: Parallel Matrix vector multiplication using MPI:

Solution file – mpiexp2.py

Steps:

- 1) The rank is identified from the mpiexec command
- 2) Inputsize/rank is the batch size allocated to each of the processes except the last one.
- 3) The last process takes the remaining data after allocation to each process before it
- 4) Matrix vector is performed in each slave
- 5) Results are returned to master

```

C:\Users\admin\Anaconda3\Scripts>mpiexec -n 4 python mpiexp2.py
rank 1
result for rank 1  [[20.09526552]
[24.24200587]
[23.01842539]
[27.43746367]
[24.14116931]
[24.76241191]
[23.77776946]
[25.11527172]
[20.17607841]
[24.30473943]
[23.08857496]
[25.75437113]
[26.83043683]
[20.4421288 ]
[23.71749182]
[22.88704724]
[25.79793987]
[23.67841687]
[24.46944637]
[24.86104392]
[22.78831934]
[24.56576931]
[24.0089882 ]
[21.92124151]
[24.94284204]]
time 0.0014758999459445477
rank 2
result for rank 2  [[25.87430532]
[22.05691752]
[20.8303714 ]
[23.47979888]
[25.50354447]

```

Core/inputsizes	10 <sup>2</sup>	10 <sup>3</sup>	10 <sup>4</sup>
1	0.0013326001353561878	0.011930100154131651	0.03754719998687506
2	0.003135299775749445	0.03129839990288019	1.193994800094515
3	0.0039114998653531075	0.04732789983972907	1.2676194002851844
4	0.015840400010347366	0.047501299995929	4.579885399900377

Exercise 3: Parallel Matrix operation using MPI:

Solution file – matrixmul.py

Steps:

- 1) The rank is identified from the mpiexec command
- 2) The matrix is broadcasted to all the processes
- 3) Inputsizes/rank is the batch size allocated to each of the processes except the last one.
- 4) The last process takes the remaining data after allocation to each process before it
- 5) Matrix multiplication is performed in each slave
- 6) Result is given to the master process by allgather() method

```

Process 0 has been given this data: [[1.83045093 1.93212788 1.51311516 ... 1.0756
02534416]
[1.33194721 1.89642837 1.05474197 ... 1.44569877 1.12979147 1.52953512]
[1.27664395 1.79011491 1.82366065 ... 1.1717437 1.6234627 1.618894 ]
...
[1.18351425 1.13784718 1.27693755 ... 1.31224477 1.51791726 1.07580222]
[1.72611064 1.08378065 1.17724872 ... 1.87944075 1.11013658 1.20362394]
[1.88091042 1.86874847 1.16528951 ... 1.3513995 1.0708804 1.09866774]]
Matrix multiplication: [[[70.99884088 70.91429066 73.22266159 ... 77.44144815 81.
74.33852777]
[73.56007732 71.92021902 73.65749405 ... 81.22422025 83.00301577
75.97964383]
[69.56673417 68.42940088 72.4513293 ... 75.79099071 78.47615037
72.6539708 ]
...
[70.24080852 69.03523395 72.15879805 ... 77.93307432 78.15465742
71.42237834]
[71.81684872 69.5068818 71.71476165 ... 78.14601128 78.07614546
72.99762927]
[73.02089048 73.03747282 74.01251403 ... 80.61253527 81.63457011
75.79343431]]]
time 1.2127240002155304

```

Core/input size	10 <sup>2</sup>	10 <sup>3</sup>	10 <sup>4</sup>
1	1. 2127240002155304	Does not run	Does not run
2	0.7399005000479519	Does not run	Does not run
3	0.5918200998567045	Does not run	Does not run
4	0.5483559998683631	Does not run	Does not run