

# Apache Spark Basics

In [ ]:

## Part a) Basic Operations on Resilient Distributed Dataset (RDD)

### 1. Perform rightOuterJoin and fullOuterJoin operations between a and b. Briefly explain your solution.

In [1114]: `import findspark  
findspark.init()  
findspark.find()  
import pyspark  
findspark.find()`

executed in 150ms, finished 07:36:01 2020-07-06

Out[1114]: 'C:\\Users\\admin\\spark\\spark-2.3.2-bin-hadoop2.7'

In [ ]:

executed in 374ms, finished 07:36:03 2020-07-06

In [ ]: `from pyspark import SparkContext, SparkConf  
from pyspark.sql import SparkSession  
conf = pyspark.SparkConf().setAppName('appName').setMaster('local')  
sc = pyspark.SparkContext(conf=conf)  
spark = SparkSession(sc)`

executed in 738ms, finished 12:23:44 2020-07-03

In [ ]:

executed in 294ms, finished 12:23:49 2020-07-03

In [1116]: `#initializing a  
a = ["spark", "rdd", "python", "context", "create", "class"]`

executed in 6ms, finished 07:36:17 2020-07-06

In [ ]:

```
In [1117]: #initializing b  
b = ["operation", "apache", "scala", "lambda","parallel","partition"]
```

executed in 5ms, finished 07:36:19 2020-07-06

```
In [ ]:
```

```
In [1122]: #creating RDDs from 'a' such that, that RDD is a key value pair  
def add_key(value):  
    return (value,1)  
rdd1 = sc.parallelize(a)  
rdd1 = rdd1.map(add_key)  
rdd1.collect()
```

executed in 728ms, finished 07:37:01 2020-07-06

```
Out[1122]: [('spark', 1),  
            ('rdd', 1),  
            ('python', 1),  
            ('context', 1),  
            ('create', 1),  
            ('class', 1)]
```

```
In [1123]: #creating RDDs from 'b' such that, that RDD is a key value pair  
rdd2 = sc.parallelize(b)  
rdd2 = rdd2.map(add_key)  
rdd2.collect()
```

executed in 1.06s, finished 07:37:03 2020-07-06

```
Out[1123]: [('operation', 1),  
            ('apache', 1),  
            ('scala', 1),  
            ('lambda', 1),  
            ('parallel', 1),  
            ('partition', 1)]
```

```
In [1124]: #Right outer join  
roj = rdd1.rightOuterJoin(rdd2)
```

executed in 36ms, finished 07:37:03 2020-07-06

```
In [1125]: roj.collect()
```

executed in 5.45s, finished 07:37:10 2020-07-06

```
Out[1125]: [('scala', (None, 1)),  
            ('parallel', (None, 1)),  
            ('partition', (None, 1)),  
            ('operation', (None, 1)),  
            ('apache', (None, 1)),  
            ('lambda', (None, 1))]
```

```
In [1126]: #Full outer join  
rdd1.fullOuterJoin(rdd2).collect()
```

executed in 5.65s, finished 07:37:17 2020-07-06

```
Out[1126]: [('python', (1, None)),  
            ('class', (1, None)),  
            ('scala', (None, 1)),  
            ('parallel', (None, 1)),  
            ('partition', (None, 1)),  
            ('spark', (1, None)),  
            ('rdd', (1, None)),  
            ('context', (1, None)),  
            ('create', (1, None)),  
            ('operation', (None, 1)),  
            ('apache', (None, 1)),  
            ('lambda', (None, 1))]
```


### Explanation

- The two inputs given are converted into two RDD's respectively such that each RDD is a key value pair with the key as the words in input and the value as 1.
- Right outer join is then computed on both RDD's
- Full outer join is also computed on both RDD's
- 1 in the value of output represents the presence of the key in the joined table.

## 2. Using map and reduce functions to count how many times the character "s" appears in all a and b

```
In [1128]: from operator import add
a = ["spark", "rdd", "python", "context", "create", "class"]
b = ["operation", "apache", "scala", "lambda", "parallel", "partition"]
#function to count and return the number of occurrences of 's' in its parameter
def detect_s(x):
    x = list(x)
    c = 0
    for i in x:
        if i == 's':
            c+=1
    return c

rdd3 = sc.parallelize(a)
rdd4 = sc.parallelize(b)
#Combining RDD's for input's a and b
newr = sc.union([rdd3,rdd4])
#Mapping Input words containing 's' to the number of 's's in that word and thereby
rdd_s = newr.map(detect_s).filter(lambda x: x is not None).reduce(add)
print(rdd_s)
```



executed in 2.56s, finished 07:43:15 2020-07-06

4

### Explanation

- Create RDD's from a and b
- Combine the RDD's
- Detect the number of 's's in a word and place it as the value in a key-value pair with the word being the key
- Reduce the values

## 3. Using aggregate function to count how many times the character "s" appears in all a and b

```
In [1142]: from operator import add
a = ["spark", "rdd", "python", "context", "create", "class"]
b = ["operation", "apache", "scala", "lambda", "parallel", "partition"]
def detect_s(x):
    x = list(x)
    c = 0
    for i in x:
        if i == 's':
            c += 1
    return c

rdd3 = sc.parallelize(a)
rdd4 = sc.parallelize(b)
newr = sc.union([rdd3, rdd4])
rdd_s = newr.filter(lambda x: 's' in x).map(detect_s).sum()
print(rdd_s)
```

executed in 1.78s, finished 08:10:53 2020-07-06

4

### Explanation

- Create RDD's from a and b
- Combine the RDD's
- Filter words with s's and pass them to a map function that maps the words with the number of 's's in that word.
- Aggregate the results using sum() aggregate function

## Part b) Basic Operations on DataFrames

```
In [1215]: #imports
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)
```

executed in 117ms, finished 09:15:29 2020-07-06

```
In [1216]: #Loading the dataset
df = sqlContext.read.json('C:\\Users\\admin\\Downloads\\students.json')
```

executed in 99ms, finished 09:15:30 2020-07-06

### 1. Replace the null value(s) in column points by the mean of all points

```
In [1217]: from pyspark.sql.functions import mean as _mean, stddev as _stddev, col
#calculating mean of points
df_stats = df.select(
    _mean(col('points')).alias("mean"),
    _stddev(col('points')).alias("std deviation")
).collect()

print(df_stats[0]["mean"])
mean = df_stats[0]['mean']
std = df_stats[0]['std deviation']
print(std)
```

executed in 103ms, finished 09:15:32 2020-07-06

11.736842105263158

3.3307007147839007

In [ ]:

executed in 6ms, finished 20:24:34 2020-07-03

```
In [1218]: #filling null values with mean of all points
df = df.na.fill(mean,['points'])
df.show()
```

executed in 150ms, finished 09:15:35 2020-07-06

course	dob	first_name	last_name	points	s_id
Humanities and Art	October 14, 1983	Alan	Joe	10	1
Computer Science	September 26, 1980	Martin	Genberg	17	2
Graphic Design	June 12, 1982	Athur	Watson	16	3
Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
Psychology	November 1, 1978	Kira	Schommer	11	5
Business	17 February 1981	Christian	Kiriam	10	6
Machine Learning	1 January 1984	Barbara	Ballard	14	7
Deep Learning	January 13, 1978	John	null	10	8
Machine Learning	26 December 1989	Marcus	Carson	15	9
Physics	30 December 1987	Marta	Brooks	11	10
Data Analytics	June 12, 1975	Holly	Schwartz	12	11
Computer Science	July 2, 1985	April	Black	11	12
Computer Science	July 22, 1980	Irene	Bradley	13	13
Psychology	7 February 1986	Mark	Weber	12	14
Informatics	May 18, 1987	Rosie	Norman	9	15
Business	August 10, 1984	Martin	Steele	7	16
Machine Learning	16 December 1990	Colin	Martinez	9	17
Data Analytics	null	Bridget	Twain	6	18
Business	7 March 1980	Darlene	Mills	19	19
Data Analytics	June 2, 1985	Zachary	null	10	20

**2. Replace the null value(s) in column dob and column last name by "unknown" and "--" respectively.**

```
In [1219]: df.show()
#filling null values with values specified in the question
df1 = df.na.fill("unknown","dob")
df1.show()
df1 = df1.na.fill("--","last_name")
df1.show()
df = df1
```

executed in 201ms, finished 09:15:37 2020-07-06

Psychology	7 February 1988	Mark	Weber	12	1
Informatics	May 18, 1987	Rosie	Norman	9	15
Business	August 10, 1984	Martin	Steele	7	16
Machine Learning	16 December 1990	Colin	Martinez	9	17
Data Analytics	unknown	Bridget	Twain	6	18
Business	7 March 1980	Darlene	Mills	19	19
Data Analytics	June 2, 1985	Zachary	null	10	20

course	dob	first_name	last_name	points	s_id
Humanities and Art	October 14, 1983	Alan	Joe	10	1
Computer Science	September 26, 1980	Martin	Genberg	17	2
Graphic Design	June 12, 1982	Athur	Watson	16	3
Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
Psychology	November 1, 1978	Kira	Schommer	11	5
Business	17 February 1981	Christian	Kiriam	10	6
Machine Learning	1 January 1984	Barbara	Ballard	14	7
Deen Learning	January 13, 1978	John	--	10	8

### Explanation for Part b) 1 and 2

- The null values in points are replaced by mean of the column points by using the mean sql function
- The null values in dob and lastname are replaced by 'unknown' and '--' respectively using fill method of pyspark

**3. In the dob column, there exist several formats of dates, e.g. October 14, 1983 and 26 December 1989. Let's convert all the dates into DD-MM-YYYY format where DD, MM and YYYY are two digits for day, two digits for months and four digits for year respectively**

**4. Insert a new column age and calculate the current age of all students**

```

In [1220]: import pyspark.sql.functions as F
from pyspark.sql.types import *
import re
#function to change the format of the date to DD-MM-YYYY
def changedateformat(value):
    #dictionary with keys as months and the values as the month's number
    months = {'January':1,'February':2,"March":3,"April":4,"May":5,"June":6,"July":7,"August":8,"September":9,"October":10,"November":11,"December":12}
    valu = str(value)
    valu1 = str(valu)
    v = valu.split()
    print(type(value))
    #regexes for detecting months, years, days from the dob column
    regmonth = r"([a-zA-Z]+)"
    regday = r"(\d+)"
    regyear = r"([0-9]{4})"

    #matching regexes with the dob column
    matchmonth = re.search(regmonth, valu)
    matchday = re.search(regday, valu)
    matchyear = re.search(regyear, valu)

    if(matchyear != None):
        year = str(matchyear.group(0))
    else:
        return "unknown"
    if(matchday!=None):
        day = str(matchday.group(0))
    else:
        return "unknown"
    #returning the result in DD_MM-YYYY format
    if(matchmonth!=None and matchday!=None and matchyear!=None):
        return day+"/"+str(months[str(matchmonth.group(0))])+"/"+year

def computeage(value):
    #matching year
    regyear = r"([0-9]{4})"
    matchyear = re.search(regyear, str(value))
    #subtracting matched year from 2020
    if(matchyear!=None):
        return 2020-int(str(matchyear.group(0)))
    return None

#convert to a UDF Function by passing in the function and return type of function
udfdate = F.udf(changedateformat, StringType())
udfage = F.udf(computeage, StringType())
df = df.withColumn("dob", udfdate("dob"))
df = df.withColumn("age",udfage("dob"))
df.show()

```

executed in 773ms, finished 09:15:40 2020-07-06

course	dob	first_name	last_name	points	s_id	age
Humanities and Art	14/10/1983	Alan	Joe	10	1	37
Computer Science	26/9/1980	Martin	Genberg	17	2	40



Graphic Design	12/6/1982	Athur	Watson	16	3	38
Graphic Design	5/4/1987	Anabelle	Sanberg	12	4	33
Psychology	1/11/1978	Kira	Schommer	11	5	42
Business	17/2/1981	Christian	Kiriam	10	6	39
Machine Learning	1/1/1984	Barbara	Ballard	14	7	36
Deep Learning	13/1/1978	John	--	10	8	42
Machine Learning	26/12/1989	Marcus	Carson	15	9	31
Physics	30/12/1987	Marta	Brooks	11	10	33
Data Analytics	12/6/1975	Holly	Schwartz	12	11	45
Computer Science	2/7/1985	April	Black	11	12	35
Computer Science	22/7/1980	Irene	Bradley	13	13	40
Psychology	7/2/1986	Mark	Weber	12	14	34
Informatics	18/5/1987	Rosie	Norman	9	15	33
Business	10/8/1984	Martin	Steele	7	16	36

### Explanation

- `changedateformat()` function takes in date of births in any format and changes it to DD/MM/YYYY format.
- `UserDefinedFunction` library is used to process the dataframe using user written functions
- `computeage()` function extracts year of birth from the dob column and calculates the age of the person if the dob is given

In [1221]: `df.show()`

executed in 663ms, finished 09:15:51 2020-07-06

course	dob	first_name	last_name	points	s_id	age
Humanities and Art	14/10/1983	Alan	Joe	10	1	37
Computer Science	26/9/1980	Martin	Genberg	17	2	40
Graphic Design	12/6/1982	Athur	Watson	16	3	38
Graphic Design	5/4/1987	Anabelle	Sanberg	12	4	33
Psychology	1/11/1978	Kira	Schommer	11	5	42
Business	17/2/1981	Christian	Kiriam	10	6	39
Machine Learning	1/1/1984	Barbara	Ballard	14	7	36
Deep Learning	13/1/1978	John	--	10	8	42
Machine Learning	26/12/1989	Marcus	Carson	15	9	31
Physics	30/12/1987	Marta	Brooks	11	10	33
Data Analytics	12/6/1975	Holly	Schwartz	12	11	45
Computer Science	2/7/1985	April	Black	11	12	35
Computer Science	22/7/1980	Irene	Bradley	13	13	40
Psychology	7/2/1986	Mark	Weber	12	14	34
Informatics	18/5/1987	Rosie	Norman	9	15	33
Business	10/8/1984	Martin	Steele	7	16	36
Machine Learning	16/12/1990	Colin	Martinez	9	17	30
Data Analytics	unknown	Bridget	Twain	6	18	null
Business	7/3/1980	Darlene	Mills	19	19	40
Data Analytics	2/6/1985	Zachary	--	10	20	35

In [510]:

executed in 105ms, finished 17:14:46 2020-07-04

## 5. Let's consider granting some points for good performed students in the class. For each student, if his point is larger than 1 standard deviation of all points, then we update his current point to 20, which is the maximum

In [1222]:

```
ints > df_stats[0]["mean"]+df_stats[0]["std deviation"] ,20).otherwise(df.points))
```

executed in 673ms, finished 09:15:55 2020-07-06

course	dob	first_name	last_name	points	s_id	age
Humanities and Art	14/10/1983	Alan	Joe	10	1	37
Computer Science	26/9/1980	Martin	Genberg	20	2	40
Graphic Design	12/6/1982	Athur	Watson	20	3	38
Graphic Design	5/4/1987	Anabelle	Sanberg	12	4	33
Psychology	1/11/1978	Kira	Schommer	11	5	42
Business	17/2/1981	Christian	Kiriam	10	6	39
Machine Learning	1/1/1984	Barbara	Ballard	14	7	36
Deep Learning	13/1/1978	John	--	10	8	42
Machine Learning	26/12/1989	Marcus	Carson	15	9	31
Physics	30/12/1987	Marta	Brooks	11	10	33
Data Analytics	12/6/1975	Holly	Schwartz	12	11	45
Computer Science	2/7/1985	April	Black	11	12	35
Computer Science	22/7/1980	Irene	Bradley	13	13	40
Psychology	7/2/1986	Mark	Weber	12	14	34
Informatics	18/5/1987	Rosie	Norman	9	15	33
Business	10/8/1984	Martin	Steele	7	16	36
Machine Learning	16/12/1990	Colin	Martinez	9	17	30
Data Analytics	unknown	Bridget	Twain	6	18	null
Business	7/3/1980	Darlene	Mills	20	19	40
Data Analytics	2/6/1985	Zachary	--	10	20	35

### Explanation

- Calculate mean+standard deviation for each point and replace points with 20 if they have more value than the sum mentioned.

## 6. Create a histogram on the new points created in the task 5

```
In [1280]: # Converting column points to list
points = df.select('points').collect()
print(type(points))
points = [int(i.points) for i in points]
print(points)
```

executed in 235ms, finished 12:10:30 2020-07-06

```
<class 'list'>
[10, 20, 20, 12, 11, 10, 14, 10, 15, 11, 12, 11, 13, 12, 9, 7, 9, 6, 20, 10]
```

```
In [ ]: #plotting a histogram
from matplotlib import pyplot as plt

plt.hist(points)
plt.xlabel('points')
plt.ylabel('frequency')
plt.title('Ex 1. Part b) 6.')
plt.show()
```

executed in 157ms, finished 12:11:02 2020-07-06

In [ ]:

## Ex 2: Manipulating Recommender Dataset with Apache Spark

### Preprocessing

```
In [1232]: #reading the dataset and preprocessing
tags = sqlContext.read.format("csv").option("delimiter",":").load('C:\\Users\\adm
tags = tags.drop('_c1').drop('_c3').drop('_c5')
```

executed in 133ms, finished 10:26:41 2020-07-06

```
In [1233]: from pyspark.sql.window import Window
```

executed in 5ms, finished 10:26:43 2020-07-06

```
In [1234]: #ordering by users and time
tags = tags.orderBy('_c0','_c6', ascending=True)
```

```
Windowspec = Window.orderBy('_c0')
#Windowspec = Windowspec.orderBy('_c6')
```

executed in 25ms, finished 10:26:45 2020-07-06

In [1235]: *#creating a new column with value of column '\_c6' from t-1*  

```
previous_time = tags.withColumn('previous_time', F.lag(tags['_c6']).over(Windowsp
```

executed in 16ms, finished 10:26:48 2020-07-06

In [1236]: `previous_time.show()`

executed in 9.58s, finished 10:26:59 2020-07-06

```
+-----+-----+-----+-----+-----+
| _c0| _c2|          _c4|          _c6|previous_time|
+-----+-----+-----+-----+-----+
| 1000| 277|children's story|1188533111|      null|
| 1000| 1994|   sci-fi. dark|1188533136| 1188533111|
| 1000| 5377|      romance|1188533150| 1188533136|
| 1000| 7147|   family bonds|1188533161| 1188533150|
| 1000| 362|animated classic|1188533171| 1188533161|
| 1000| 276|      family|1188533235| 1188533171|
|10003|42013|    Passable|1150432435| 1188533235|
|10003|51662|  FIOS on demand|1207953326| 1150432435|
|10003|54997|  FIOS on demand|1207953335| 1207953326|
|10003|55765|  FIOS on demand|1207953342| 1207953335|
|10003|55363|  FIOS on demand|1207953420| 1207953342|
|10003|56152|  FIOS on demand|1207953526| 1207953420|
|10003|55116|  FIOS on demand|1207953636| 1207953526|
|10003|56174|  FIOS on demand|1207953670| 1207953636|
|10003|55176|  FIOS on demand|1207953755| 1207953670|
|10003|55247|  FIOS on demand|1207953756| 1207953755|
|10003|54881|  FIOS on demand|1207953758| 1207953756|
|10003|55820|  FIOS on demand|1207953873| 1207953758|
|10003|53123|  FIOS on demand|1207953875| 1207953873|
|10003|53550|  FIOS on demand|1207953937| 1207953875|
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

In [1237]: *#calculating time difference between values in column \_c6 at time t-1 and t*  

```
result = previous_time.withColumn('time_difference', (previous_time['_c6']-previo
```

executed in 10ms, finished 10:27:18 2020-07-06

In [1238]: `result.show()`

executed in 7.82s, finished 10:27:27 2020-07-06

_c0	_c2	_c4	_c6	previous_time	time_difference
1000	277	children's story	1188533111	null	null
1000	1994	sci-fi. dark	1188533136	1188533111	25.0
1000	5377	romance	1188533150	1188533136	14.0
1000	7147	family bonds	1188533161	1188533150	11.0
1000	362	animated classic	1188533171	1188533161	10.0
1000	276	family	1188533235	1188533171	64.0
10003	42013	Passable	1150432435	1188533235	-3.81008E7
10003	51662	FIOS on demand	1207953326	1150432435	5.7520891E7
10003	54997	FIOS on demand	1207953335	1207953326	9.0
10003	55765	FIOS on demand	1207953342	1207953335	7.0
10003	55363	FIOS on demand	1207953420	1207953342	78.0
10003	56152	FIOS on demand	1207953526	1207953420	106.0
10003	55116	FIOS on demand	1207953636	1207953526	110.0
10003	56174	FIOS on demand	1207953670	1207953636	34.0
10003	55176	FIOS on demand	1207953755	1207953670	85.0
10003	55247	FIOS on demand	1207953756	1207953755	1.0
10003	54881	FIOS on demand	1207953758	1207953756	2.0
10003	55820	FIOS on demand	1207953873	1207953758	115.0
10003	53123	FIOS on demand	1207953875	1207953873	2.0
10003	53550	FIOS on demand	1207953937	1207953875	62.0

only showing top 20 rows

```
In [1240]: prev_time= 0
prev_user = 0
session = 0
#assigns session based on time difference
def assign_session(time, user):
    global prev_time
    global prev_user
    global session
    if(user!=prev_user):
        time = float("-inf")
    if(time== None or time > 1800 or user!=prev_user or time< -1800):
        session+=1
    prev_user = user
    return session
udfsession = F.udf(assign_session, StringType())

tags_session = result.withColumn("session",udfsession("time_difference",'_c0'))
```

executed in 35ms, finished 10:36:33 2020-07-06

## 1. Tagging session for each user

In [1241]: tags\_session.show()

executed in 9.58s, finished 10:36:45 2020-07-06

_c0	_c2	_c4	_c6	previous_time	time_difference	session
1000	277	children's story	1188533111	null	null	1
1000	1994	sci-fi. dark	1188533136	1188533111	25.0	1
1000	5377	romance	1188533150	1188533136	14.0	1
1000	7147	family bonds	1188533161	1188533150	11.0	1
1000	362	animated classic	1188533171	1188533161	10.0	1
1000	276	family	1188533235	1188533171	64.0	1
10003	42013	Passable	1150432435	1188533235	-3.81008E7	2
10003	51662	FIOS on demand	1207953326	1150432435	5.7520891E7	3
10003	54997	FIOS on demand	1207953335	1207953326	9.0	3
10003	55765	FIOS on demand	1207953342	1207953335	7.0	3
10003	55363	FIOS on demand	1207953420	1207953342	78.0	3
10003	56152	FIOS on demand	1207953526	1207953420	106.0	3
10003	55116	FIOS on demand	1207953636	1207953526	110.0	3
10003	56174	FIOS on demand	1207953670	1207953636	34.0	3
10003	55176	FIOS on demand	1207953755	1207953670	85.0	3
10003	55247	FIOS on demand	1207953756	1207953755	1.0	3
10003	54881	FIOS on demand	1207953758	1207953756	2.0	3
10003	55820	FIOS on demand	1207953873	1207953758	115.0	3
10003	53123	FIOS on demand	1207953875	1207953873	2.0	3
10003	53550	FIOS on demand	1207953937	1207953875	62.0	3

only showing top 20 rows

### Explanation

- Preprocess the dataset to have an additional column that describes the time difference between two tags of a user
- Assign a new session if the time difference is greater than 30 minutes = 1800 seconds.

## 2.Frequency of tagging for each user session

In [1242]: freq\_tag = tags\_session.orderBy('session').groupBy('session').count()

executed in 21ms, finished 10:37:46 2020-07-06

```
In [1243]: freq_tag.show()
```

```
executed in 9.65s, finished 10:37:56 2020-07-06
```

```
+-----+-----+
|session|count|
+-----+-----+
|      1|    6|
|     10|    1|
|    100|    1|
|   1000|    6|
|  10000|    5|
| 100001|    8|
| 100002|    1|
| 100003|    1|
| 100004|    7|
| 100005|    3|
| 100006|    2|
| 100007|    2|
| 100008|    1|
| 100009|    1|
|   1001|    1|
|  10010|    1|
|  10011|    2|
|  10012|    1|
|  10013|    2|
|  10014|    1|
+-----+-----+
```

only showing top 20 rows

### Explanation

- Perform groupBy on session attribute and count() such that the frequency of tagging within each user session is obtained.

## 3. Mean and standard deviation of the tagging frequency for each user

```
In [1255]: per_user = tags_session.groupBy('_c0','session').count().orderBy('_c0','session')
per_user.show()
```

executed in 7.60s, finished 11:13:26 2020-07-06

```
+-----+-----+-----+
|  _c0|session|count|
+-----+-----+-----+
| 1000|      1|     6|
|10003|      2|     1|
|10003|      3|    18|
|10003|      4|    38|
|10020|      5|     2|
|10025|      6|     1|
|10032|     10|     1|
|10032|     11|     4|
|10032|     12|     1|
|10032|     13|     1|
|10032|     14|     4|
|10032|     15|     1|
|10032|     16|     1|
|10032|     17|     1|
|10032|     18|     1|
|10032|      7|    39|
|10032|      8|     1|
|10032|      9|     1|
|10051|     19|     1|
|10058|     20|    35|
+-----+-----+-----+
only showing top 20 rows
```

- The above table shows the number of tags performed for each session by each user.



```
In [1269]: #mean of tagging frequency of each user
per_user_mean = per_user.groupBy('_c0').mean().orderBy('_c0')
per_user_mean.show()
```

executed in 2m 2s, finished 11:50:39 2020-07-06

```
+-----+-----+
|_c0|      avg(count)|
+-----+-----+
| 1000|              6.0|
|10003|             19.0|
|10020|              2.0|
|10025|              1.0|
|10032| 4.66666666666667|
|10051|              1.0|
|10058|25.333333333333332|
|10059|              2.5|
|10064|              1.0|
|10084|             3.75|
|10094|              2.0|
| 1010|              4.0|
|10117|              1.5|
|10125|             21.0|
|10132|          1.5625|
|10154|              8.0|
|10167|              1.0|
| 1017|              7.0|
|10181|             11.0|
|10191|2.666666666666665|
+-----+-----+
only showing top 20 rows
```

### Explanation

- This table calculates the mean of tagging frequency of each user by performing groupBy on user and then calculating the mean

```
In [1266]: #std of tagging frequency of each user
per_user.groupBy('_c0').agg({'count':'std'}).show()
```

executed in 2m 26s, finished 11:35:05 2020-07-06

```
+-----+-----+
|  _c0|      stddev(count)|
+-----+-----+
|11563|                NaN|
| 1436|                NaN|
|17427|0.7071067811865476|
| 2136|                NaN|
|23318|                NaN|
|28473|                0.0|
| 2904|                NaN|
|29549|                NaN|
|32812|1.4142135623730951|
|38271|                NaN|
|39917|                NaN|
|42688|                NaN|
|44446|                NaN|
|48370|                0.0|
| 5325|                NaN|
|57051|                NaN|
|57113|                NaN|
```

### Explanation

- This table calculates the std of tagging frequency of each user by performing groupBy on user and then calculating the mean

## 4. Mean and standard deviation of the tagging frequency across users

```
In [1270]: freq_stats = freq_tag.select(
            _mean(col('count')).alias("mean"),
            _stddev(col('count')).alias("std deviation")
        ).collect()
```

```
freq_mean = (freq_stats[0]["mean"])
freq_std = (freq_stats[0]["std deviation"])
```

executed in 12.4s, finished 11:50:51 2020-07-06

```
In [1271]: print(freq_mean, freq_std)
```

executed in 6ms, finished 11:51:12 2020-07-06

```
6.269596589045589 20.519115697996106
```

## 5. Provide the list of users with a mean tagging frequency within the two standard deviation from the mean frequency of all users.

```
In [1277]: m = per_user_mean.select('_c0').filter(per_user_mean['avg(count)']<= freq_mean+2*stddev_samp(freq_mean))
```



executed in 225ms, finished 12:04:13 2020-07-06

```
In [1278]: m.show()
```

executed in 2m 9s, finished 12:06:23 2020-07-06

```
+-----+
|  _c0 |
+-----+
| 1000 |
|10003 |
|10020 |
|10025 |
|10032 |
|10051 |
|10058 |
|10059 |
|10064 |
|10084 |
|10094 |
| 1010 |
|10117 |
|10125 |
|10132 |
|10154 |
|10167 |
| 1017 |
|10181 |
|10191 |
+-----+
```

only showing top 20 rows

## Explanation

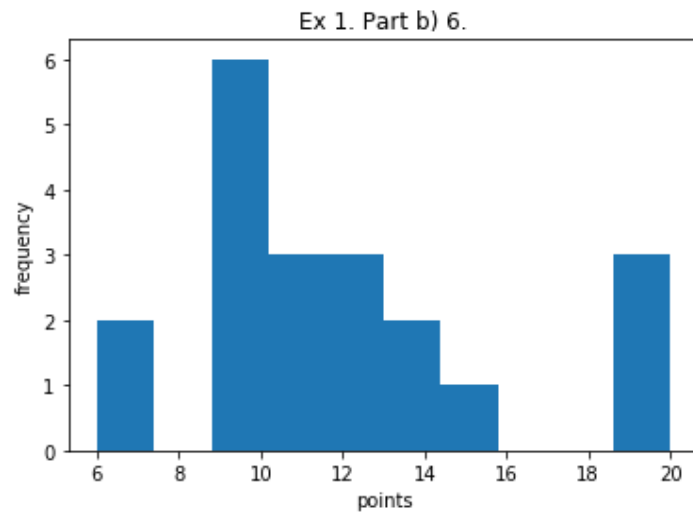
- Filtering users whose per user mean is within two stds from the mean across all users

## Solution to Exercise sheet 8

### Exercise 1:

#### Part b)

#### 6. Histogram



#### Snapshot of final dataset:

course	dob	first_name	last_name	points	s_id	age
Humanities and Art	14/10/1983	Alan	Joe	10	1	37
Computer Science	26/9/1980	Martin	Genberg	20	2	40
Graphic Design	12/6/1982	Athur	Watson	20	3	38
Graphic Design	5/4/1987	Anabelle	Sanberg	12	4	33
Psychology	1/11/1978	Kira	Schommer	11	5	42
Business	17/2/1981	Christian	Kiriam	10	6	39
Machine Learning	1/1/1984	Barbara	Ballard	14	7	36
Deep Learning	13/1/1978	John	--	10	8	42
Machine Learning	26/12/1989	Marcus	Carson	15	9	31
Physics	30/12/1987	Marta	Brooks	11	10	33
Data Analytics	12/6/1975	Holly	Schwartz	12	11	45
Computer Science	2/7/1985	April	Black	11	12	35
Computer Science	22/7/1980	Irene	Bradley	13	13	40
Psychology	7/2/1986	Mark	Weber	12	14	34
Informatics	18/5/1987	Rosie	Norman	9	15	33
Business	10/8/1984	Martin	Steele	7	16	36
Machine Learning	16/12/1990	Colin	Martinez	9	17	30
Data Analytics	unknown	Bridget	Twain	6	18	null
Business	7/3/1980	Darlene	Mills	20	19	40
Data Analytics	2/6/1985	Zachary	--	10	20	35